

# Physical Design



CHIDATA

# Physical Design

---

How do we organize a “real” collection of data

**Last few lectures:** Single data type, simple list

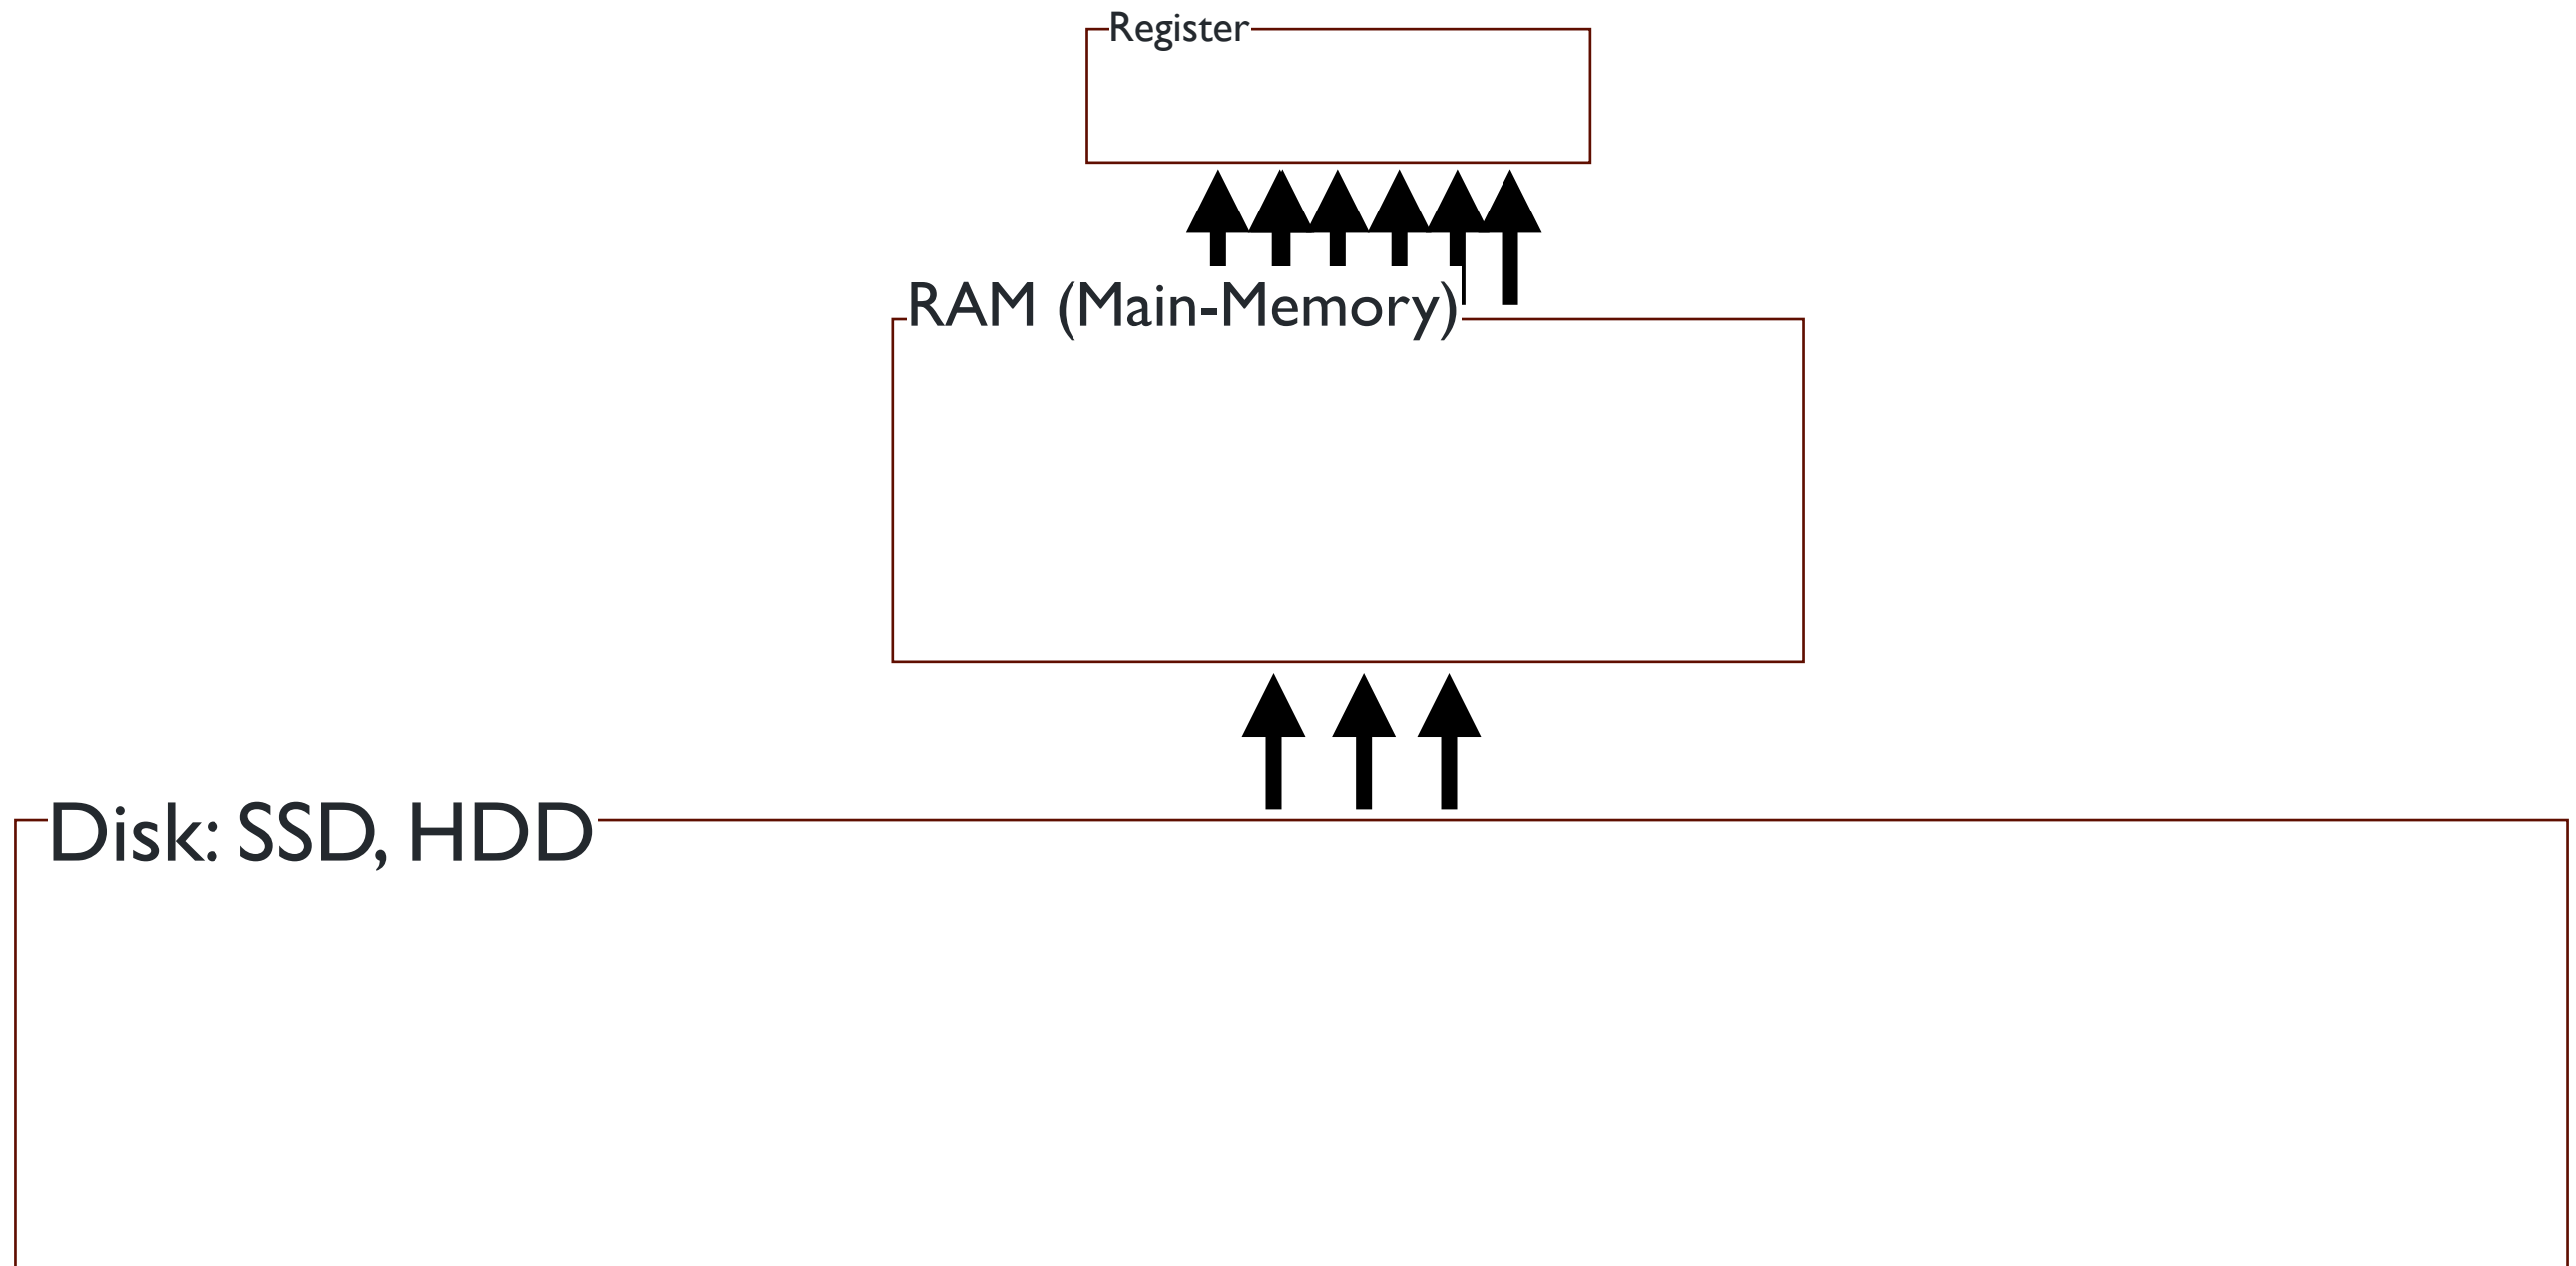
```
data = ['red', 'blue', 'red', 'red']
```

**This lecture:** Collection of “records”

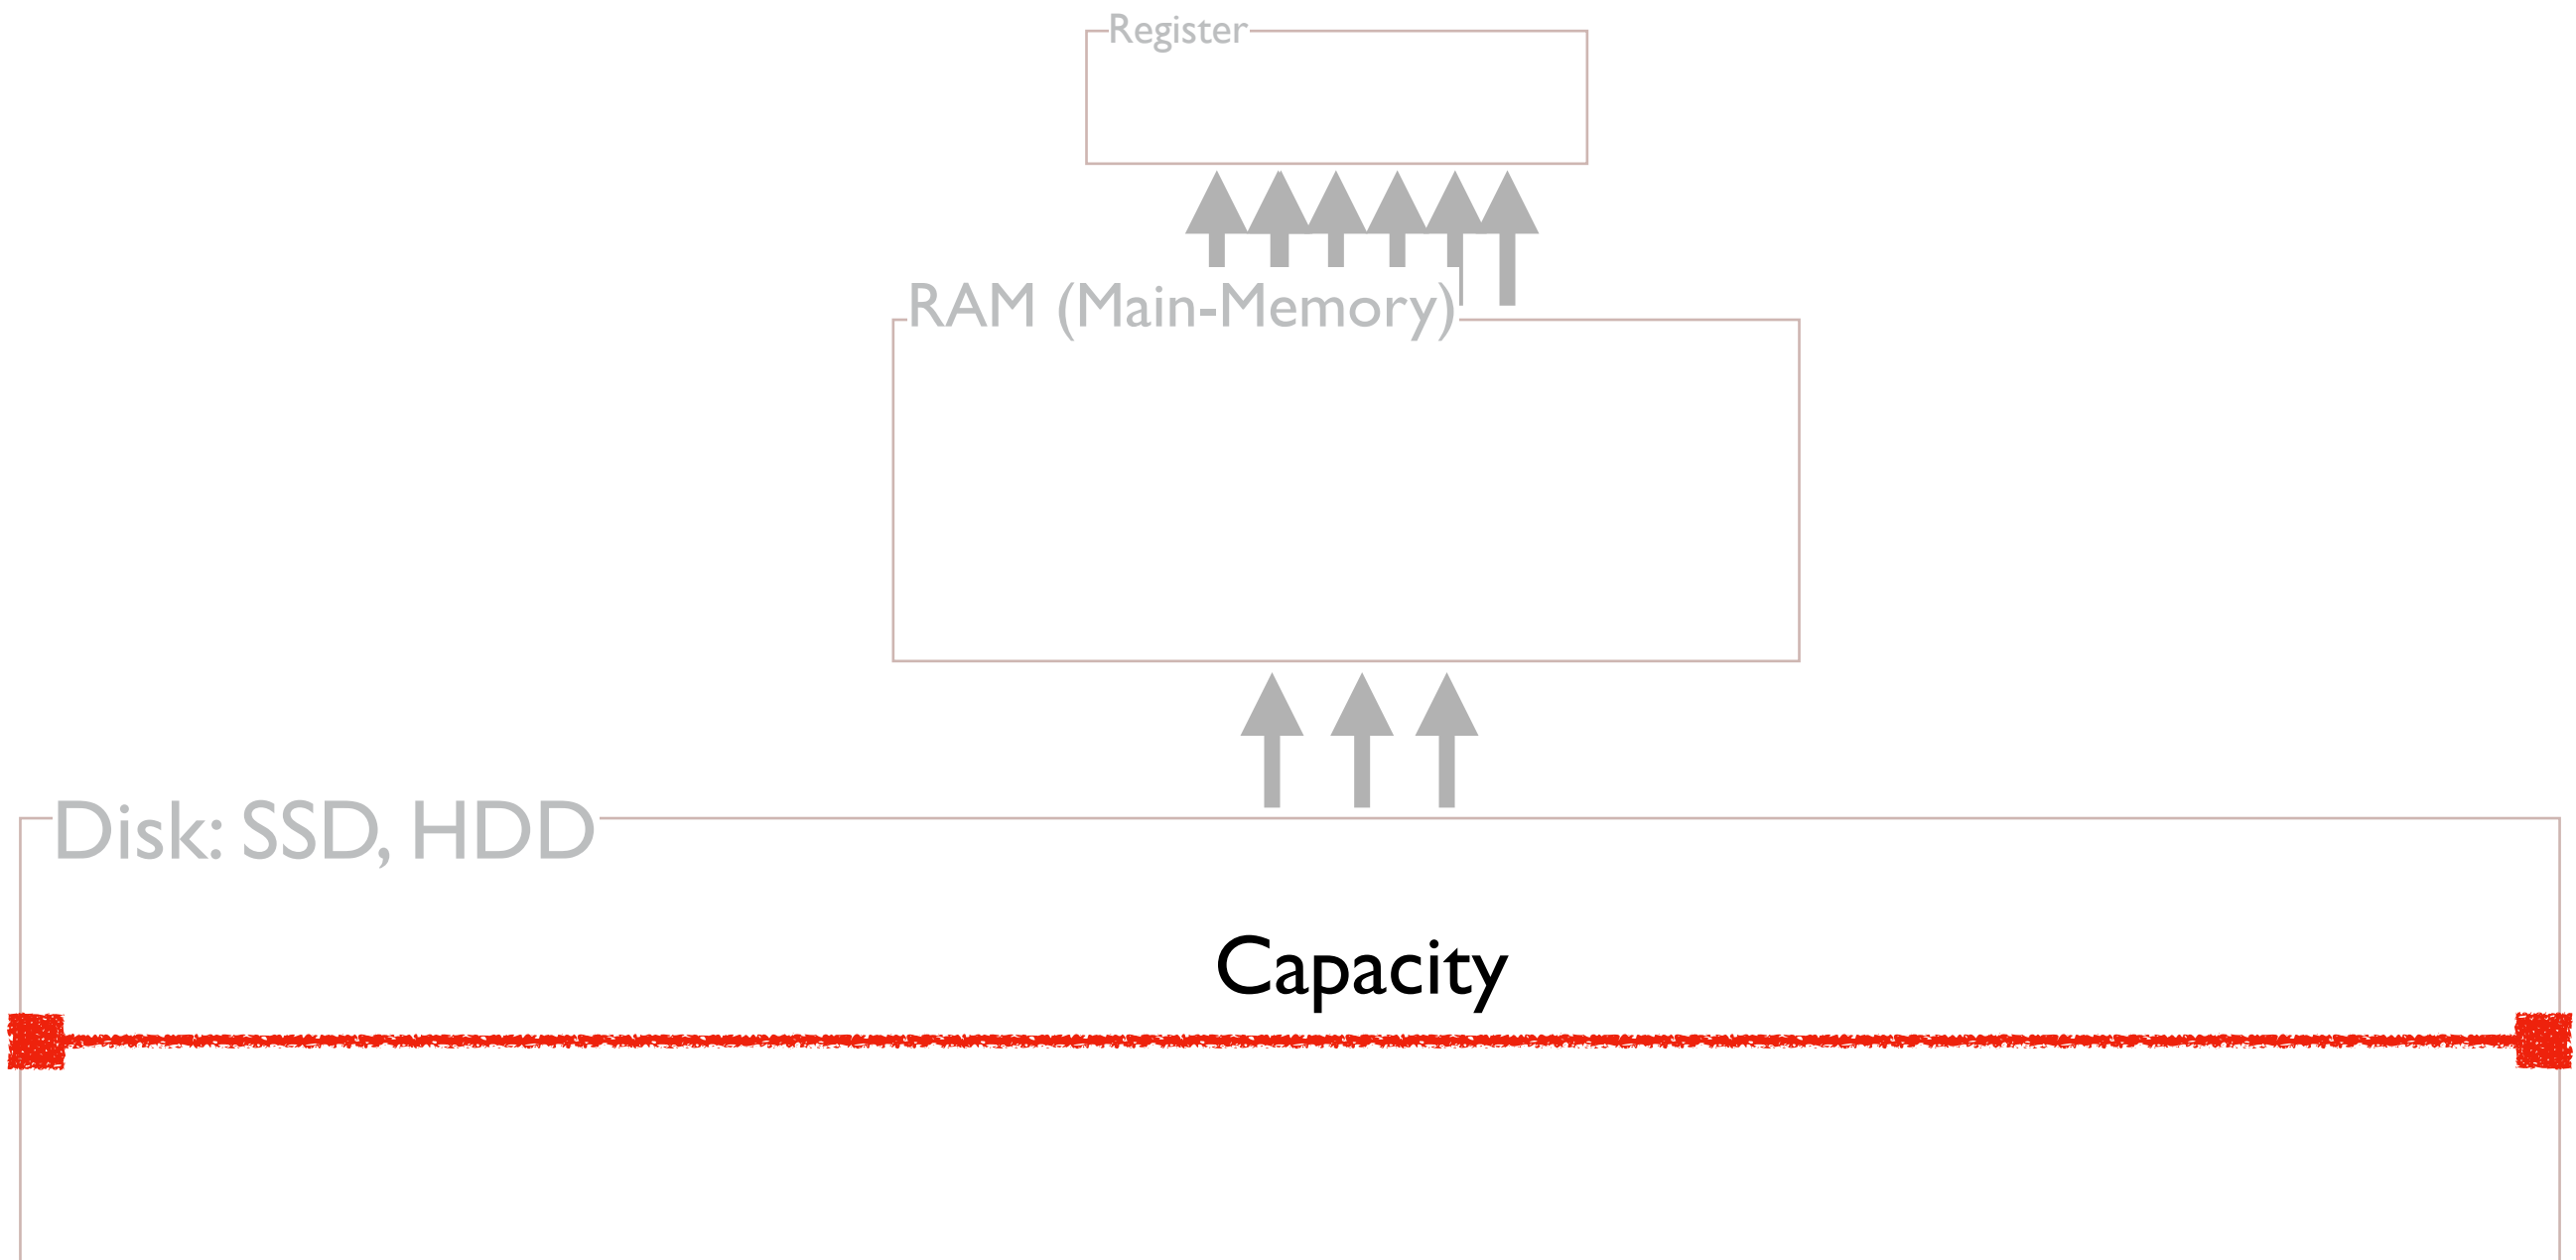
Car_id	Make	Model
1	Toyota	Camry
2	Ford	Fusion
3	Toyota	Corolla
4	Honda	Civic

# Schematic of Computer Storage

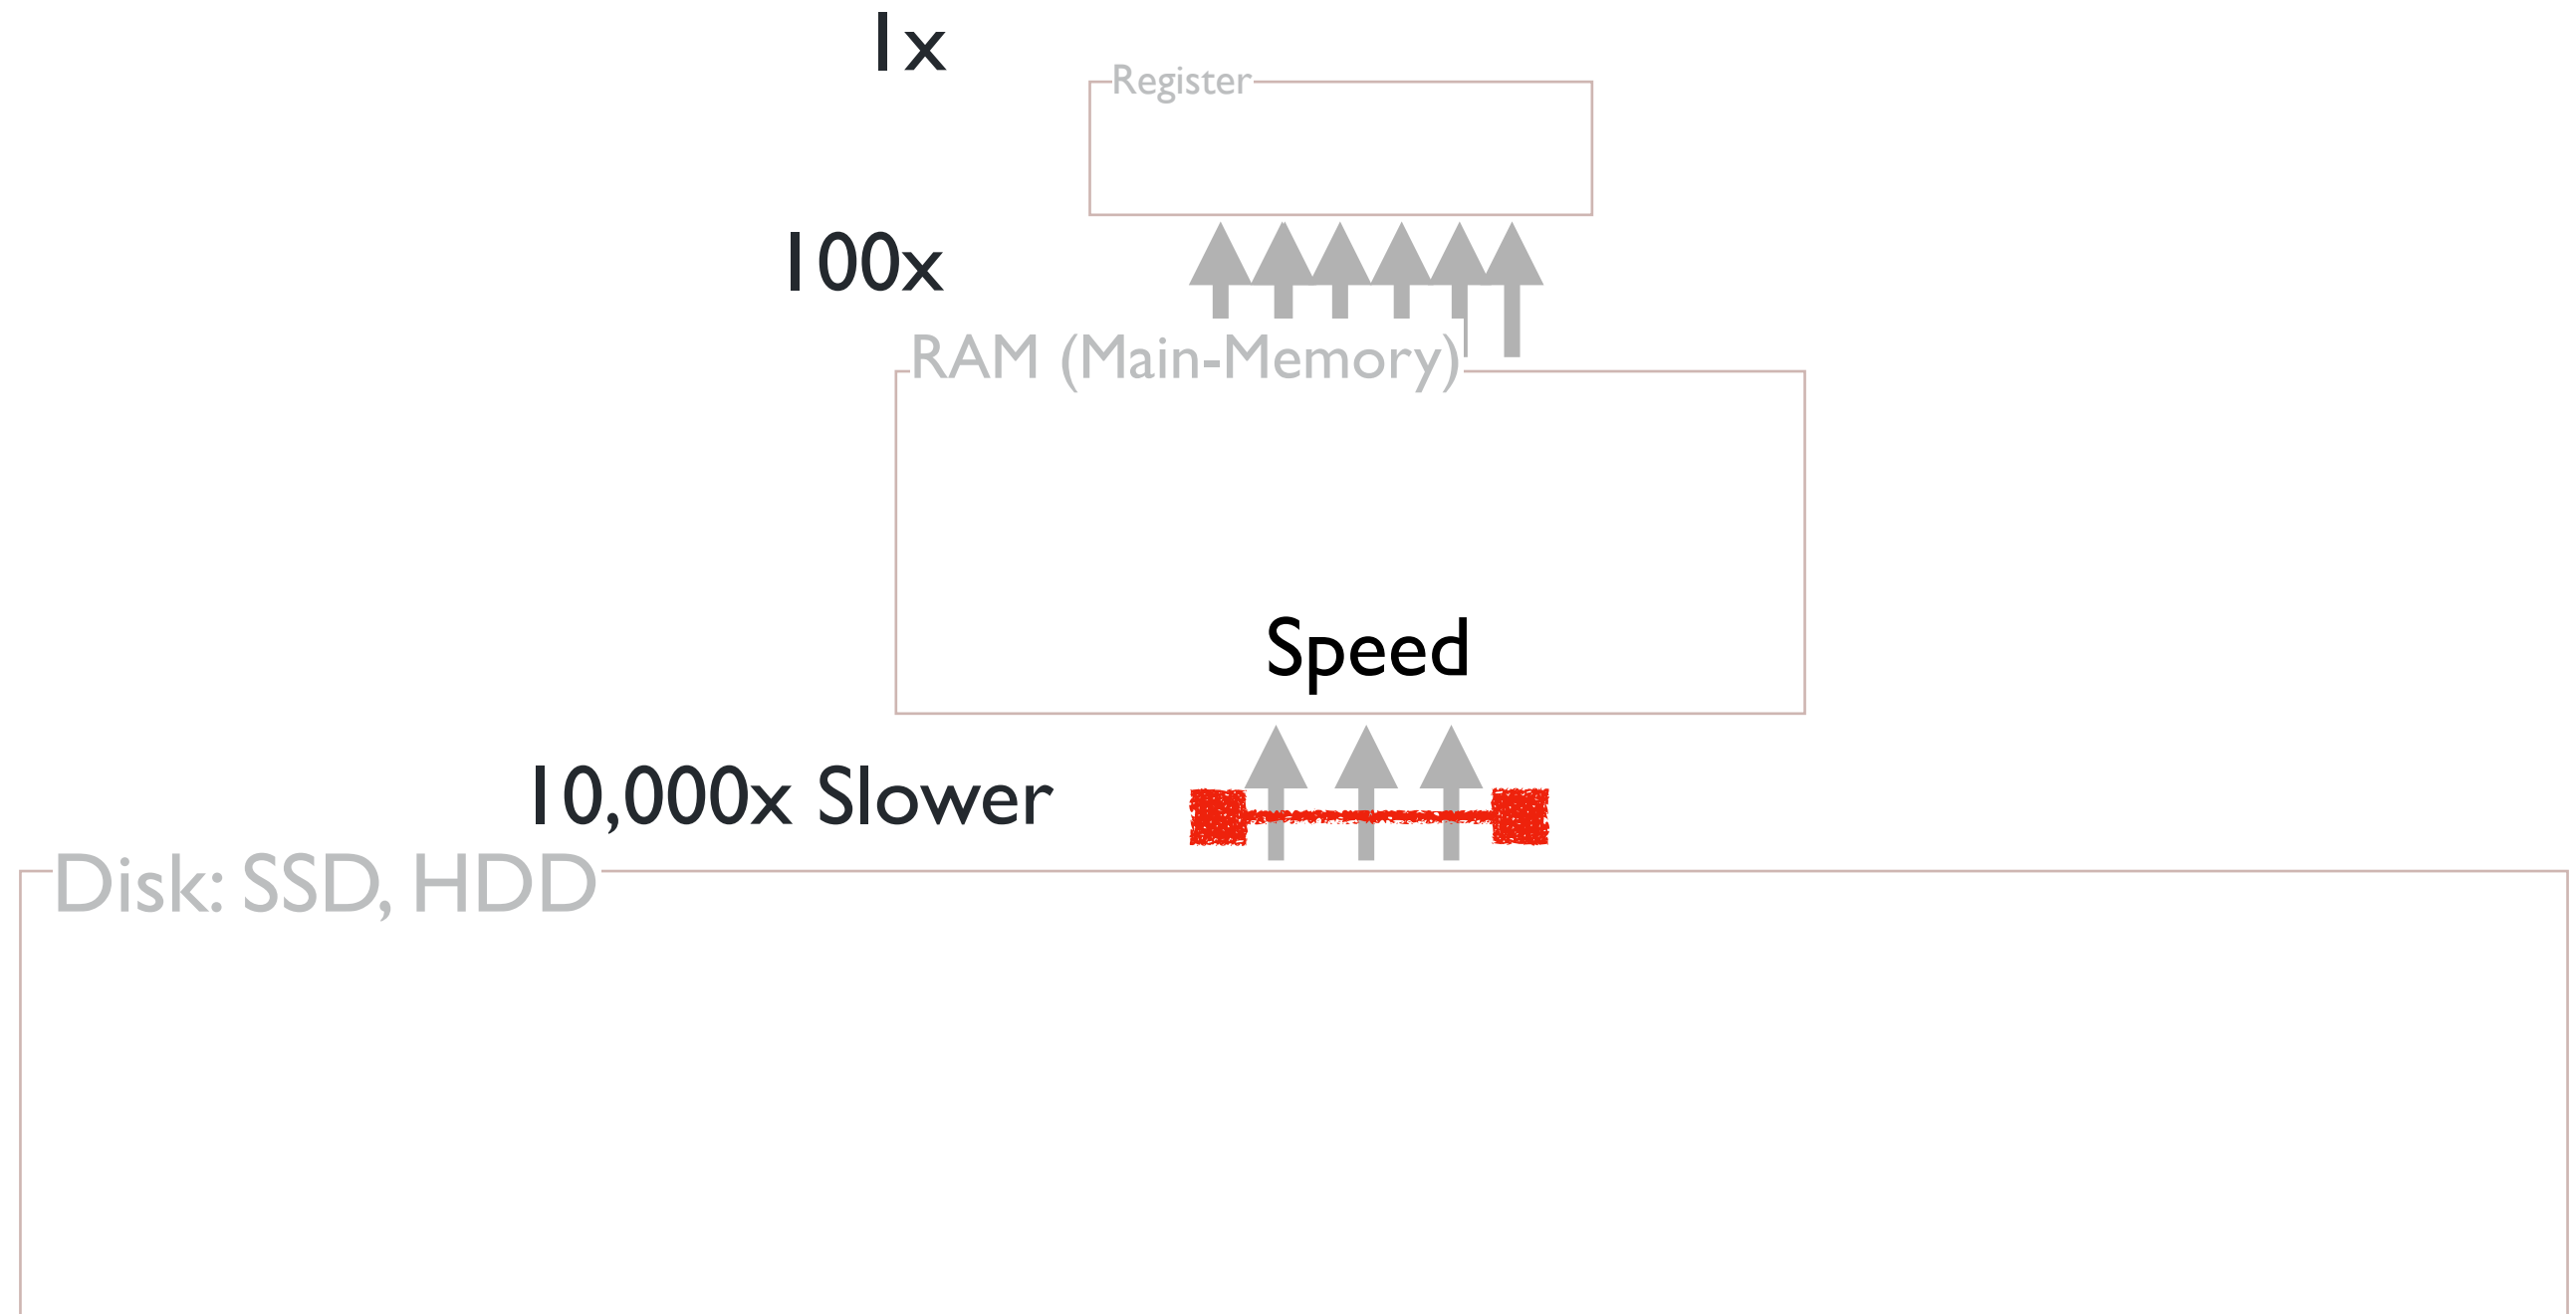
---




# Schematic of Computer Storage



# Schematic of Computer Storage




# Data Must Move For Analysis



# Physical Design Objective

---

Organize data so you avoid moving data that you don't need.



Think about data as stored in blocks

# Row-Oriented Storage

Store data as a collection of “records”

Car_id	Make	Model
1	Toyota	Camry
2	Ford	Fusion
3	Toyota	Corolla
4	Honda	Civic

Disk: SSD, HDD

**Block 1**

Car_id	Make	Model
1	Toyota	Camry
2	Ford	Fusion

*Min: 1, Max: 2*

**Block 2**

Car_id	Make	Model
3	Toyota	Corolla
4	Honda	Civic

*Min: 3, Max: 4*



# Row-Oriented Storage

Task: read record with id 2

I/O Cost: **Size of a block**

Disk: SSD, HDD

**Block 1**

Car_id	Make	Model
1	Toyota	Camry
2	Ford	Fusion

*Min: 1, Max: 2*

**Block 2**

Car_id	Make	Model
3	Toyota	Corolla
4	Honda	Civic

*Min: 3, Max: 4*

# Row-Oriented Storage

Task: count records with make = 'Toyota'

I/O Cost: **Size of all blocks**

Disk: SSD, HDD

**Block 1**

Car_id	Make	Model
1	Toyota	Camry
2	Ford	Fusion

**Block 2**

Car_id	Make	Model
3	Toyota	Corolla
4	Honda	Civic

# Row-Oriented Storage

---

Task: read record with id 2

I/O Cost: **Size of a block**

Pretty much as good as it gets

Task: count records with make = 'Toyota'

I/O Cost: **Size of all blocks**

Hmm.....

Sort the data on “Make”

Car_id	Make	Model
2	Ford	Fusion
4	Honda	Civic
1	Toyota	Camry
3	Toyota	Corolla

Disk: SSD, HDD

Block 1

Car_id	Make	Model
2	Ford	Fusion
4	Honda	Civic

Min: Ford, Max: Honda

Block 2

Car_id	Make	Model
1	Toyota	Camry
3	Toyota	Corolla

Min: Toyota, Max: Toyota

# Let's try that again

Task: count records with make = 'Toyota'

Car_id	Make	Model
2	Ford	Fusion
4	Honda	Civic
1	Toyota	Camry
3	Toyota	Corolla

Disk: SSD, HDD

**Block 1**

Car_id	Make	Model
2	Ford	Fusion
4	Honda	Civic

*Min: Ford, Max: Honda*

**Block 2**

Car_id	Make	Model
1	Toyota	Camry
3	Toyota	Corolla

*Min: Toyota, Max: Toyota*

# Uh-oh!

Task: read record with id 2

Car_id	Make	Model
2	Ford	Fusion
4	Honda	Civic
1	Toyota	Camry
3	Toyota	Corolla

Disk: SSD, HDD

**Block 1**

Car_id	Make	Model
2	Ford	Fusion
4	Honda	Civic

*Min: Ford, Max: Honda*

**Block 2**

Car_id	Make	Model
1	Toyota	Camry
3	Toyota	Corolla

*Min: Toyota, Max: Toyota*

# Sort on “Make”

---

Task: read record with id 2

I/O Cost: **Size of all blocks (expected to see half blocks)**

Hmm....

Task: count records with make = ‘Toyota’

I/O Cost: **Size of all blocks that contain Toyota**

Pretty much as good as it gets

# Row-Oriented Layouts

---

Blocks are collections of roughly the same amount of records

Sorting allows us to filter on particular attributes: need to select these before hand.



# Insertions

Sort the data on “Make”

Car_id	Make	Model
2	Ford	Fusion
4	Honda	Civic
1	Toyota	Camry
3	Toyota	Corolla
5	Ford	Escape

Disk: SSD, HDD

**Block 1**

Car_id	Make	Model
2	Ford	Fusion
4	Honda	Civic

*Min: Ford, Max: Honda*

**Block 2**

Car_id	Make	Model
1	Toyota	Camry
3	Toyota	Corolla

*Min: Toyota, Max: Toyota*

# Insertions

Sort the data on “Make”

Car_id	Make	Model
2	Ford	Fusion
4	Honda	Civic
1	Toyota	Camry
3	Toyota	Corolla
5	Ford	Escape

Disk: SSD, HDD

**Block 1**

Car_id	Make	Model
2	Ford	Fusion
5	Ford	Fusion

Min: Ford, Max: Ford

**Block 2**

Car_id	Make	Model
4	Honda	Civic
1	Toyota	Camry

Min: Honda, Max: Toyota

**Block 3**

Car_id	Make	Model
3	Toyota	Corolla

Min: Toyota, Max: Toyota

# Row-Oriented Layouts

---

Blocks are collections of roughly the same amount of records

Sorting allows us to filter on particular attributes: need to select these before hand.

Sorting: insertions are bad\*

\* appends are ok (know all new records are  $\geq$  max val).

# An Alternate Approach?

Task: count records with make = 'Toyota'

Car_id	Make	Model
1	Toyota	Camry
2	Ford	Fusion
3	Toyota	Corolla
4	Honda	Civic

Disk: SSD, HDD

**Block 1**

Car_id	Make	Model
1	Toyota	Camry
2	Ford	Fusion

**Block 2**

Car_id	Make	Model
3	Toyota	Corolla
4	Honda	Civic

# An Alternate Approach?

Task: count records with make = 'Toyota'

Disk: SSD, HDD

**Block 1**

Car_id	Make	Model
1	Toyota	Camry
2	Ford	Fusion

**Block 2**

Car_id	Make	Model
3	Toyota	Corolla
4	Honda	Civic

Note: Each block contains 4 strings

# Columnar Storage

Task: count records with make = 'Toyota'

Car_id	Make	Model
1	Toyota	Camry
2	Ford	Fusion
3	Toyota	Corolla
4	Honda	Civic

Disk: SSD, HDD

**Block 1**

Make
Toyota
Ford
Toyota
Honda

**Block 2**

Model
Camry
Fusion
Corolla
Civic

# Columnar Storage

Task: count records with make = 'Toyota'

I/O Cost: **Size of all “Make” blocks**

Disk: SSD, HDD

**Block 1**

Make
Toyota
Ford
Toyota
Honda

**Block 2**

Model
Camry
Fusion
Corolla
Civic

Can dictionary encode and compress each block to fit more!

# Agressive Compression

## Dictionary Encoding

Disk: SSD, HDD

**Block 1**

Make
"00"
"01"
"00"
"11"

**Block 2**

Model
"00"
"01"
"10"
"11"

~~Task: count records with make = 'Toyota'~~

Task: count records with make = '00'

Can fit much more relevant data (effectively) in one block.



# Columnar Storage

Task: read record with id 2

I/O Cost: **Size of all blocks**

Disk: SSD, HDD

**Block 1**

Make
Toyota
Ford
Toyota
Honda

**Block 2**

Model
Camry
Fusion
Corolla
Civic

Makes sense when you are interested in slicing or aggregating along columns!

# Column-Oriented Layouts

---

Blocks are collections of columns of data


Efficient aggregation along certain columns

Enables aggressive compression (similar to our previous lectures)

# Physical Design Objective

---

Organize data so you avoid moving data that you don't need.



Disk: SSD, HDD

Many modern systems use a hybrid of row and columnar storage

# Evaluation Metrics

---

**Workload:** Which filters are fast? Which are slow?

**Storage size:** How big is the stored data?

**Maintenance:** How much effort does it take to support insertions or updates to the data?